

HW 3: Priority-based Scheduler for xv6

Note: The pictures were taken after I worked on all the task, I will make sure take pictures as I go on next time

Task 1. Modify the provided ps command to print the priority of each process.

When defining system calls, the first thing you have to do is add them to syscall.c, users.h, and syscall.h to define them. We then create our functions in the sysproc.c, getpriority is the simpler function that just returns the priority from my proc. And for the setpriority we initialize a priority int, from there we see if there is space for the priority and if there is we set the priority attribute on the proc struct to the size of that address and return 0, if it does not have any space return -1. Because procinfo returns the value of all the information of the given processes, we will initialize the priority attribute in this function so we can add it to the information that it returns.

```
$ QEMU: Terminated
david321239@ubuntu1:~/Desktop/0s2$ make qemu
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp
1 -nographic -drive file=fs.img,if=none,format=raw,id=x0 -device virtio-blk-devi
ce,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

init: starting sh
$ ps
pid      state      size      ppid      name      priority  cputime  age
1        sleeping   12288     0         init       0         0
2        sleeping   16384     1         sh         0         0
3        running    12288     2         ps         0         0       20
```

Task 2. Add a readytime field to struct proc, initialize it correctly, and modify ps to print a process's age.

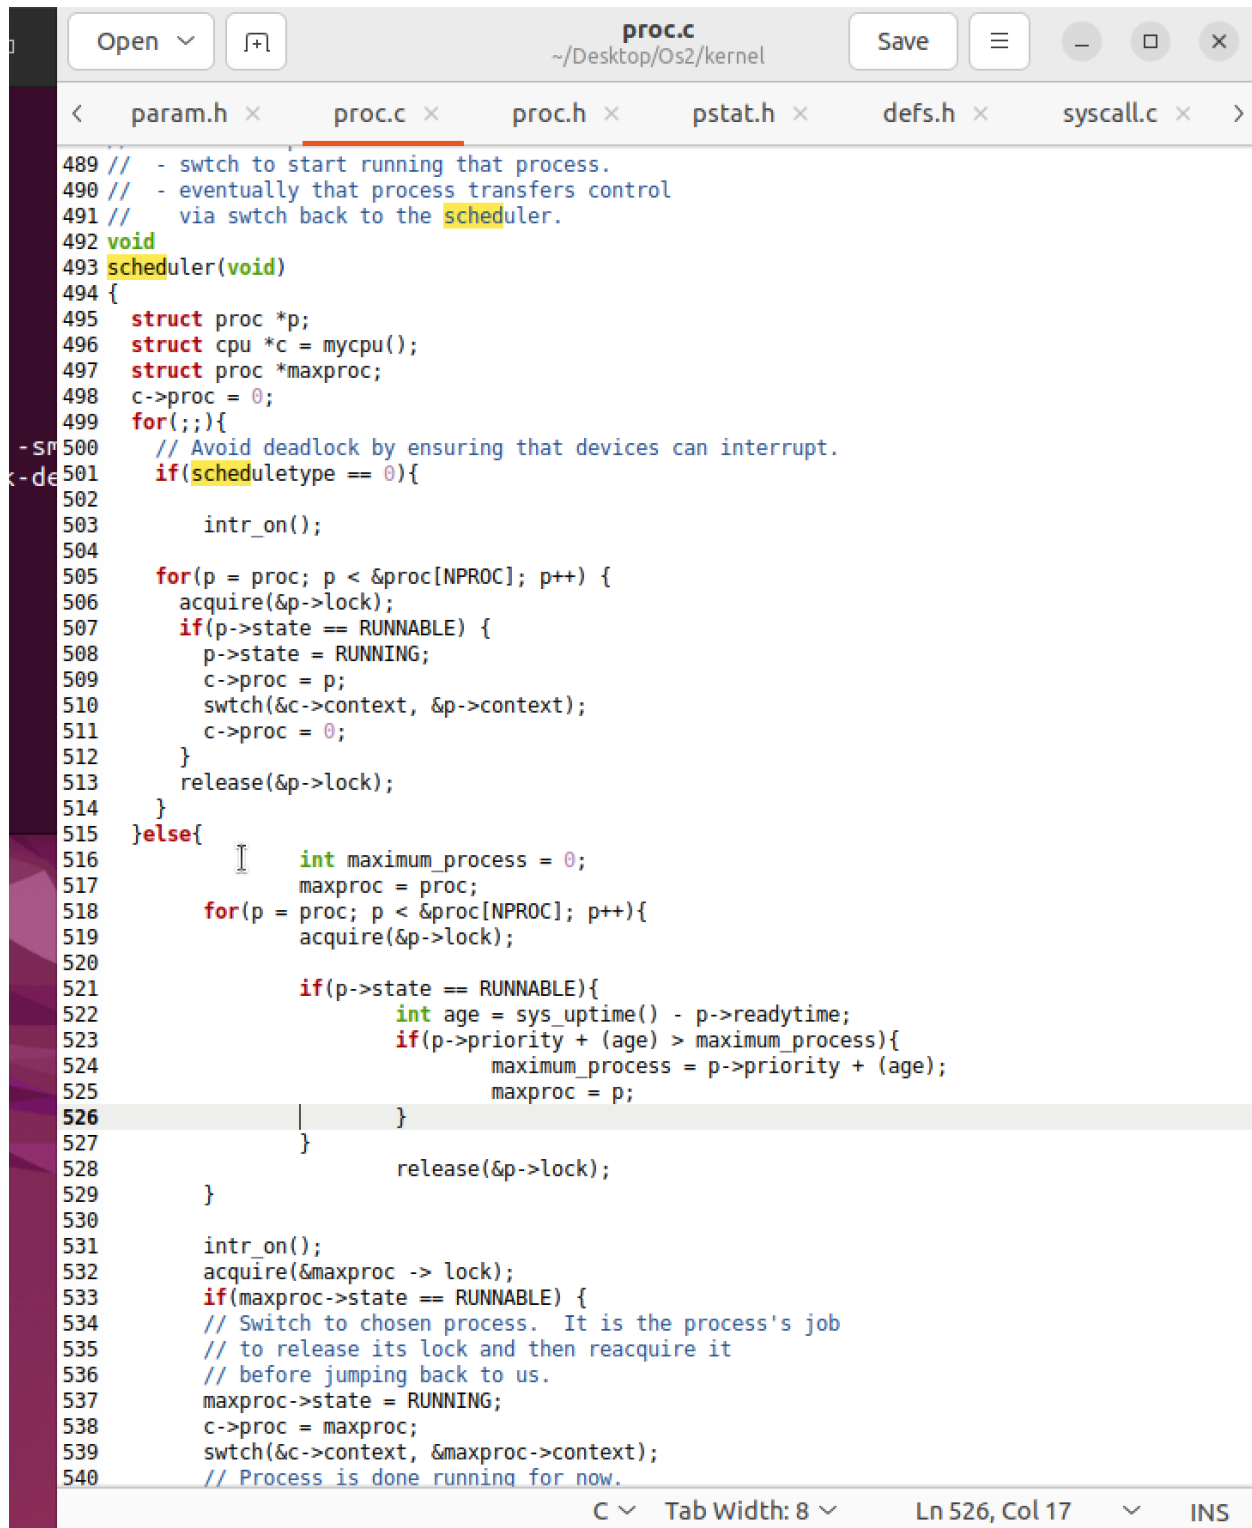
The process has an attribute called readytime. It will initialize when the process is changed to running then in ps.c to get the age by getting the current time and subtracting by the amount of time it was in the runnable state(readytime).

```
$ QEMU: Terminated
david321239@ubuntu1:~/Desktop/0s2$ make qemu
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp
1 -nographic -drive file=fs.img,if=none,format=raw,id=x0 -device virtio-blk-devi
ce,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

init: starting sh
$ ps
pid      state     size      ppid      name      priority  cputime  age
1        sleeping  12288      0         init       0          0
2        sleeping  16384      1         sh         0          0
3        running   12288      2         ps         0          0      20
$
```

Task 3. Implement a priority-based scheduler.



```
489 // - switch to start running that process.
490 // - eventually that process transfers control
491 //   via switch back to the scheduler.
492 void
493 scheduler(void)
494 {
495     struct proc *p;
496     struct cpu *c = mycpu();
497     struct proc *maxproc;
498     c->proc = 0;
499     for(;;){
500         // Avoid deadlock by ensuring that devices can interrupt.
501         if(schedulertype == 0){
502             intr_on();
503
504             for(p = proc; p < &proc[NPROC]; p++) {
505                 acquire(&p->lock);
506                 if(p->state == RUNNABLE) {
507                     p->state = RUNNING;
508                     c->proc = p;
509                     swtch(&c->context, &p->context);
510                     c->proc = 0;
511                 }
512                 release(&p->lock);
513             }
514         }else{
515             int maximum_process = 0;
516             maxproc = proc;
517             for(p = proc; p < &proc[NPROC]; p++){
518                 acquire(&p->lock);
519
520                 if(p->state == RUNNABLE){
521                     int age = sys_uptime() - p->readytime;
522                     if(p->priority + (age) > maximum_process){
523                         maximum_process = p->priority + (age);
524                         maxproc = p;
525                     }
526                 }
527             }
528             release(&p->lock);
529         }
530
531         intr_on();
532         acquire(&maxproc->lock);
533         if(maxproc->state == RUNNABLE) {
534             // Switch to chosen process. It is the process's job
535             // to release its lock and then reacquire it
536             // before jumping back to us.
537             maxproc->state = RUNNING;
538             c->proc = maxproc;
539             swtch(&c->context, &maxproc->context);
540             // Process is done running for now.
```

We change the scheduler method in proc.c and we add the scheduling types in param.h. In proc.c we had edited the scheduler method to include a priority based scheduler. We will now have two types of schedulers priority and roundrobin. If the type is 0 it is going to be round robin otherwise it will priority-based scheduling. I have learned how to implement a scheduler

into my operating system codebase by making modifications to the proc file and adding scheduling types in the param.h header. A challenge I faced was in the actual implementation of the priority scheduler as well as some annoying debugging.

Task 4. Add aging to your priority-based scheduler.

Our policy is that if our process plus age is greater than maximum processes, which is declared as a local variable that holds the process plus its ageing, will define if that processes will be running next. I learned how to calculate age then using the processes ages and in this case using it in a priority scheduler. Some difficulties I ran into was navigating through the xv6 directories to make sure that all of the attributes were properly passing their values.

```
david321239@ubuntu1:~/Desktop/0s2$ make qemu
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp
1 -nographic -drive file=fs.img,if=none,format=raw,id=x0 -device virtio-blk-devi
ce,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

init: starting sh
$ ps
pid      state      size      ppid      name      priority  cputime  age
1        sleeping   12288     0         init      0         0
2        sleeping   16384     1         sh         0         0
3        running    12288     2         ps         0         0        20
$ pexec 5 matmul 5&; matmul 10 &
$ pexec 10 ps
pid      state      size      ppid      name      priority  cputime  age
1        sleeping   12288     0         init      0         0
2        sleeping   16384     1         sh         0         0
8        runnable   12288     6         matmul     0         0
7        runnable   12288     1         matmul     0         0
6        sleeping   12288     1         pexec      0         0
9        sleeping   12288     2         pexec      0         0
10       running    12288     9         ps         0         0        24366
$ Time: 60 ticks
Time: 122 ticks
```